



CS1699: Blockchain Technology and Cryptocurrency

8. Mechanics of Bitcoin

Bill Laboon

Bitcoin Consensus

- ❖ Append-only ledger (via blockchain)
- ❖ Decentralized consensus (via consensus mechanisms)
- ❖ Validation of transactions (via miners & proof-of-work)

Account-Based Ledger: Fungible Coins - Signatories And Account Balances

BLOCK	ACTION	SIGNATORY
1	CREATE 25 COINS; TRANSFER ALL TO ALICE	MINER
2	TRANSFER 10 COINS TO BOB	ALICE
3	TRANSFER 5 COINS TO CAROL	ALICE
4	TRANSFER 2 COINS TO DAN	BOB

Account-Based Ledger

- ❖ Accounts (could still be delineated by public keys) would have a certain number of bitcoins
- ❖ A transaction would e.g. take five coins from Alice and give five coins to Bob

Account-Based Ledger: Efficiency Issues

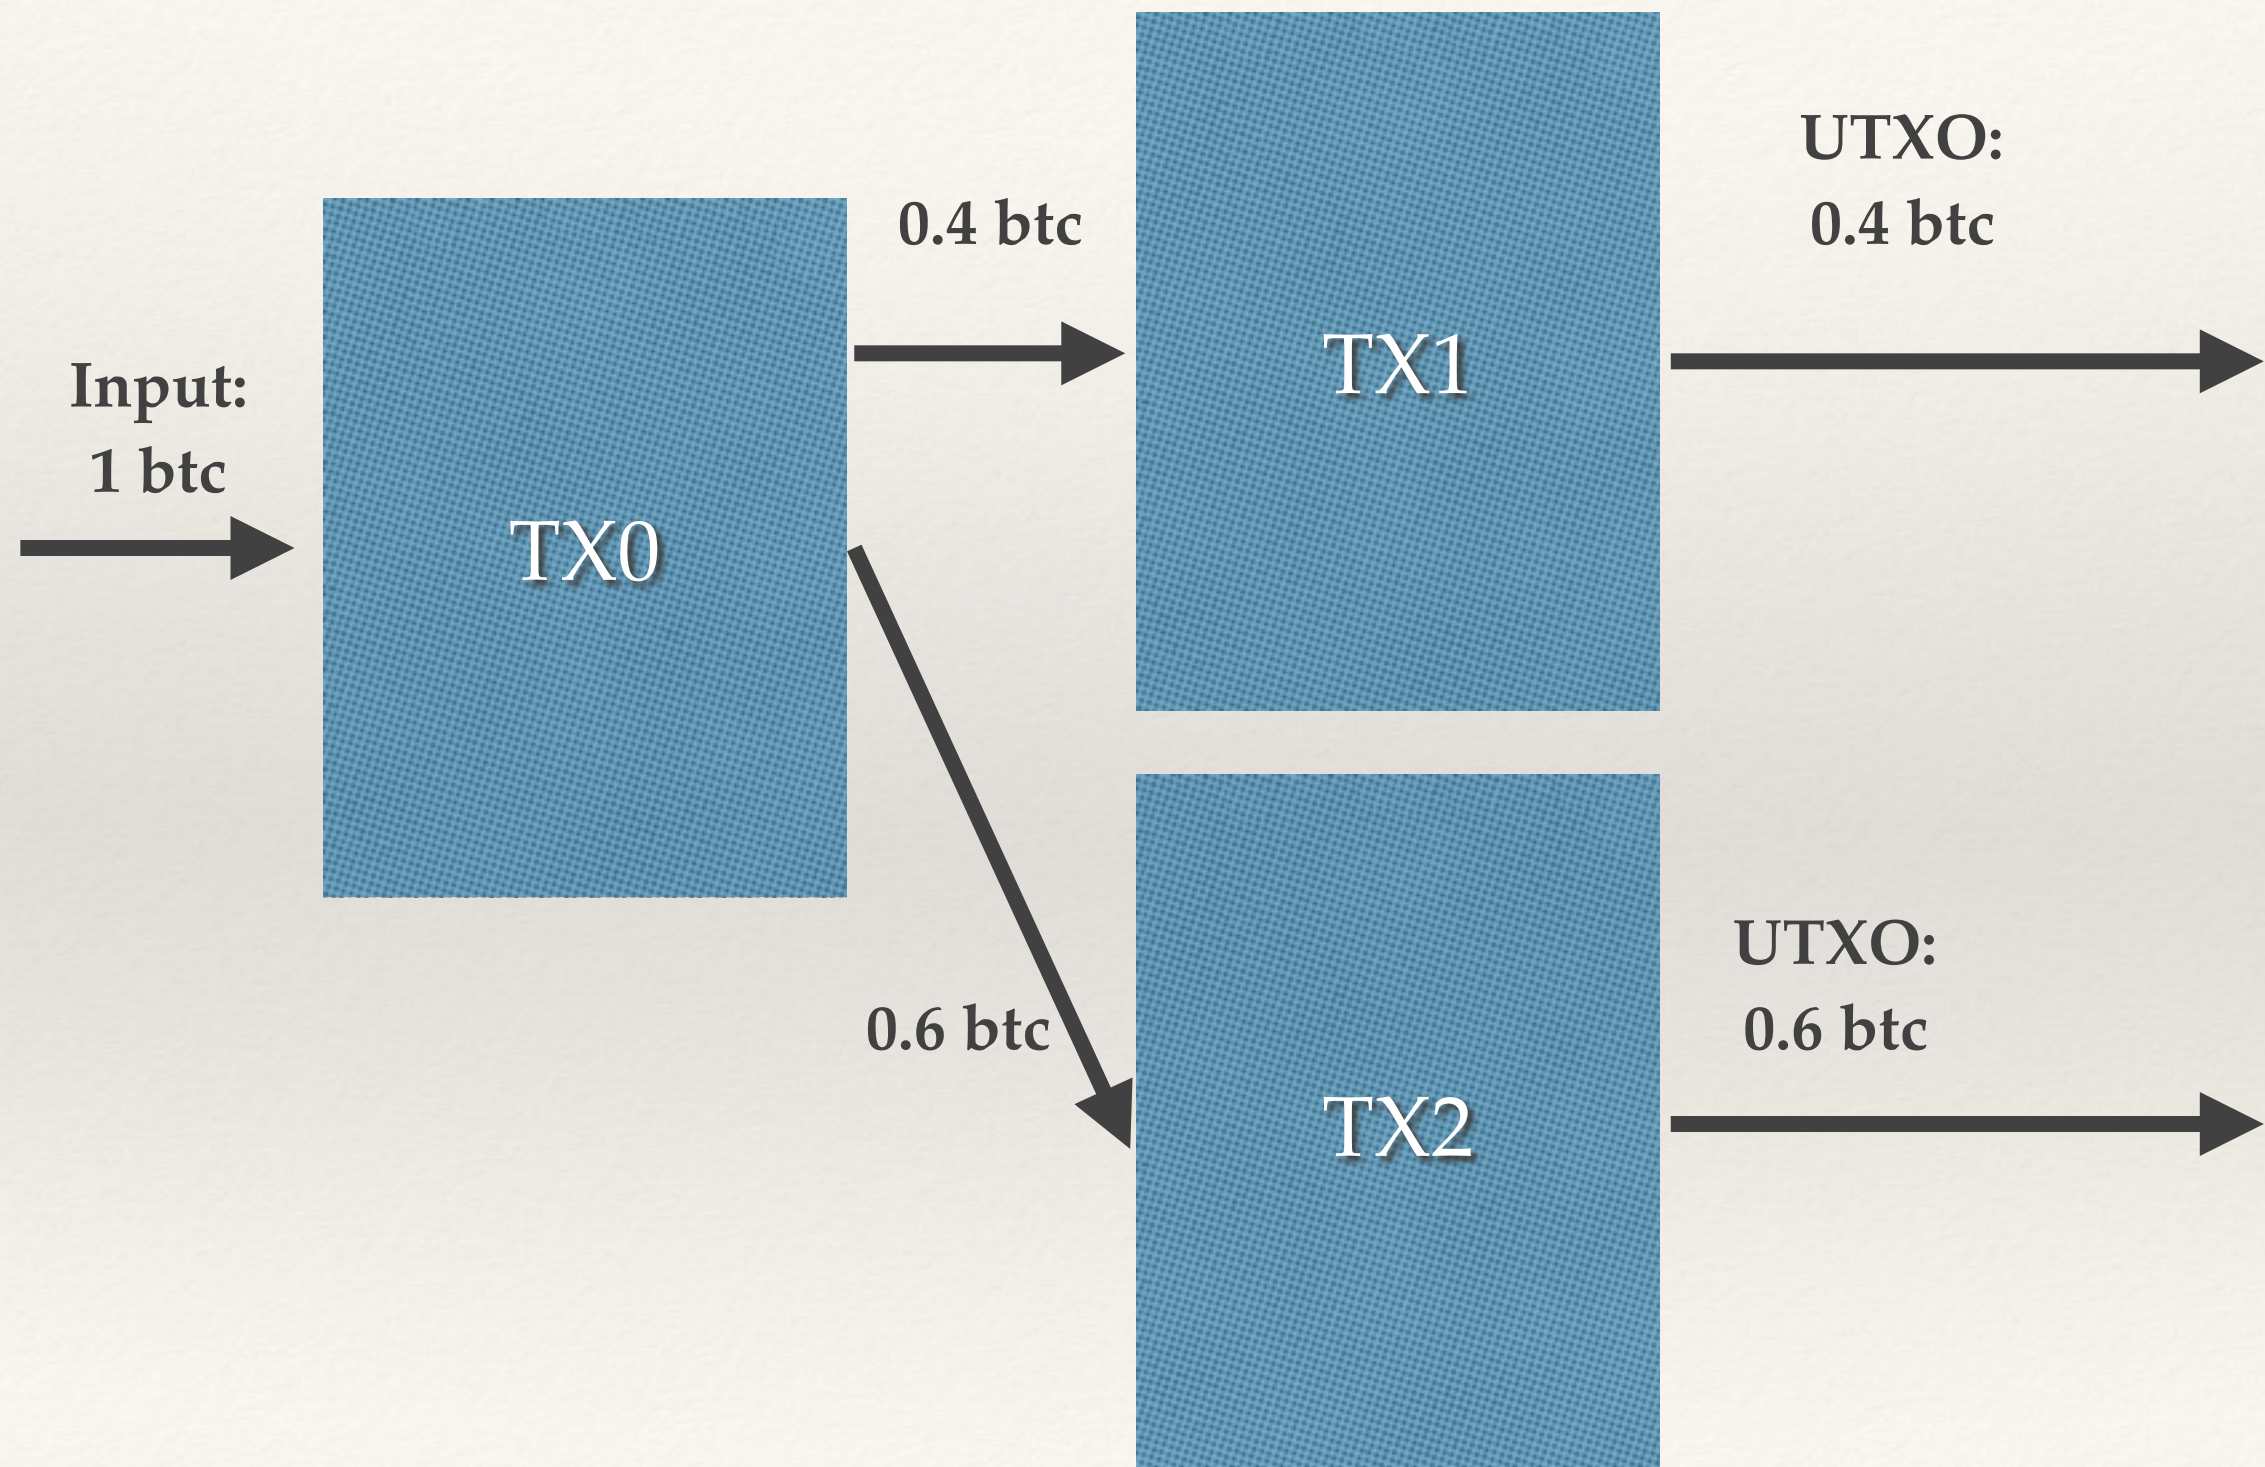
Q: Is the transaction in red valid? A: May have to track back to genesis block!

BLOCK	ACTION	SIGNATORY
1	CREATE 25 COINS; TRANSFER ALL TO ALICE	MINER
2	TRANSFER 10 COINS TO BOB	ALICE
3	TRANSFER 5 COINS TO CAROL	ALICE
4	TRANSFER 2 COINS TO DAN	BOB

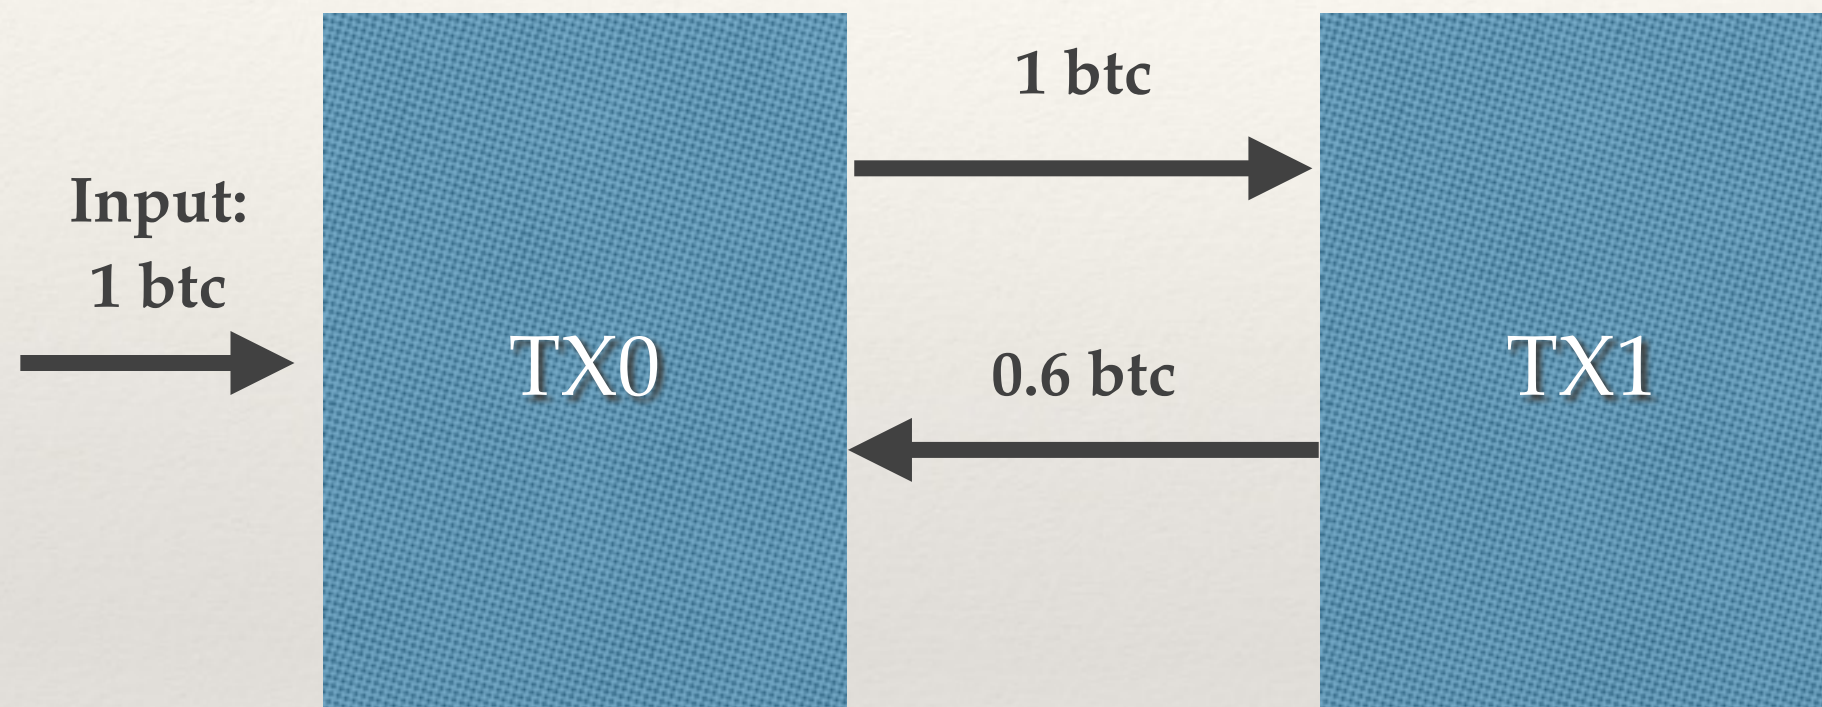
Transaction-Based Model (Bitcoin)

- ❖ Every transaction has *valid* inputs and outputs
 - ❖ Properly signed
 - ❖ Inputs must be consumed entirely
 - ❖ Total value of outputs < inputs
- ❖ “Accounts” are just a collection of *UTXOs* (*unspent transaction outputs*)

Transactions Behind the Scenes

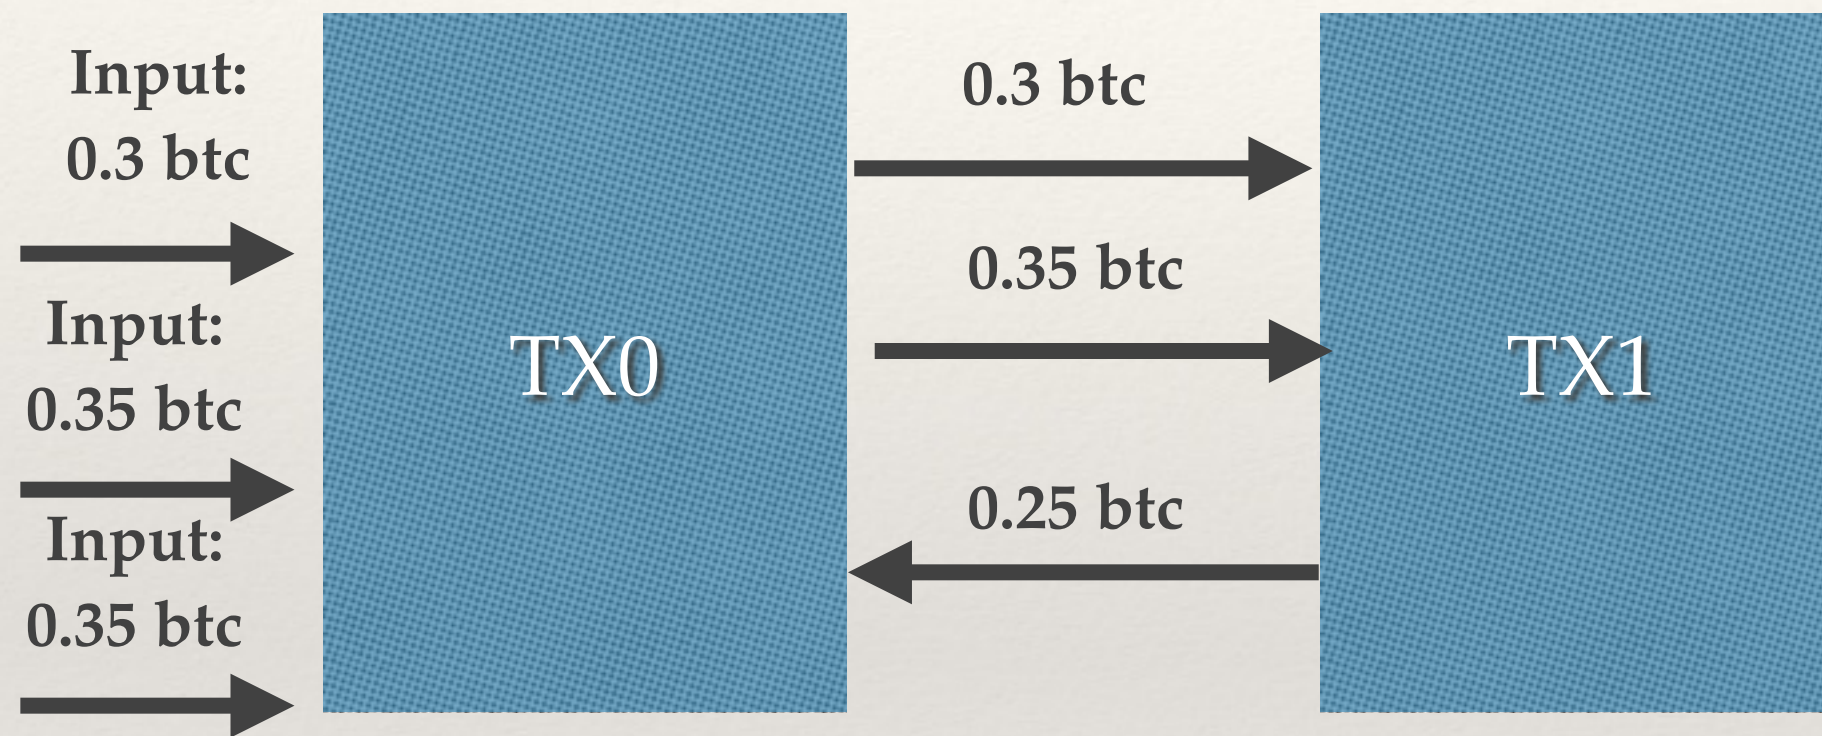


Splitting TXs - Sending 0.4 btc



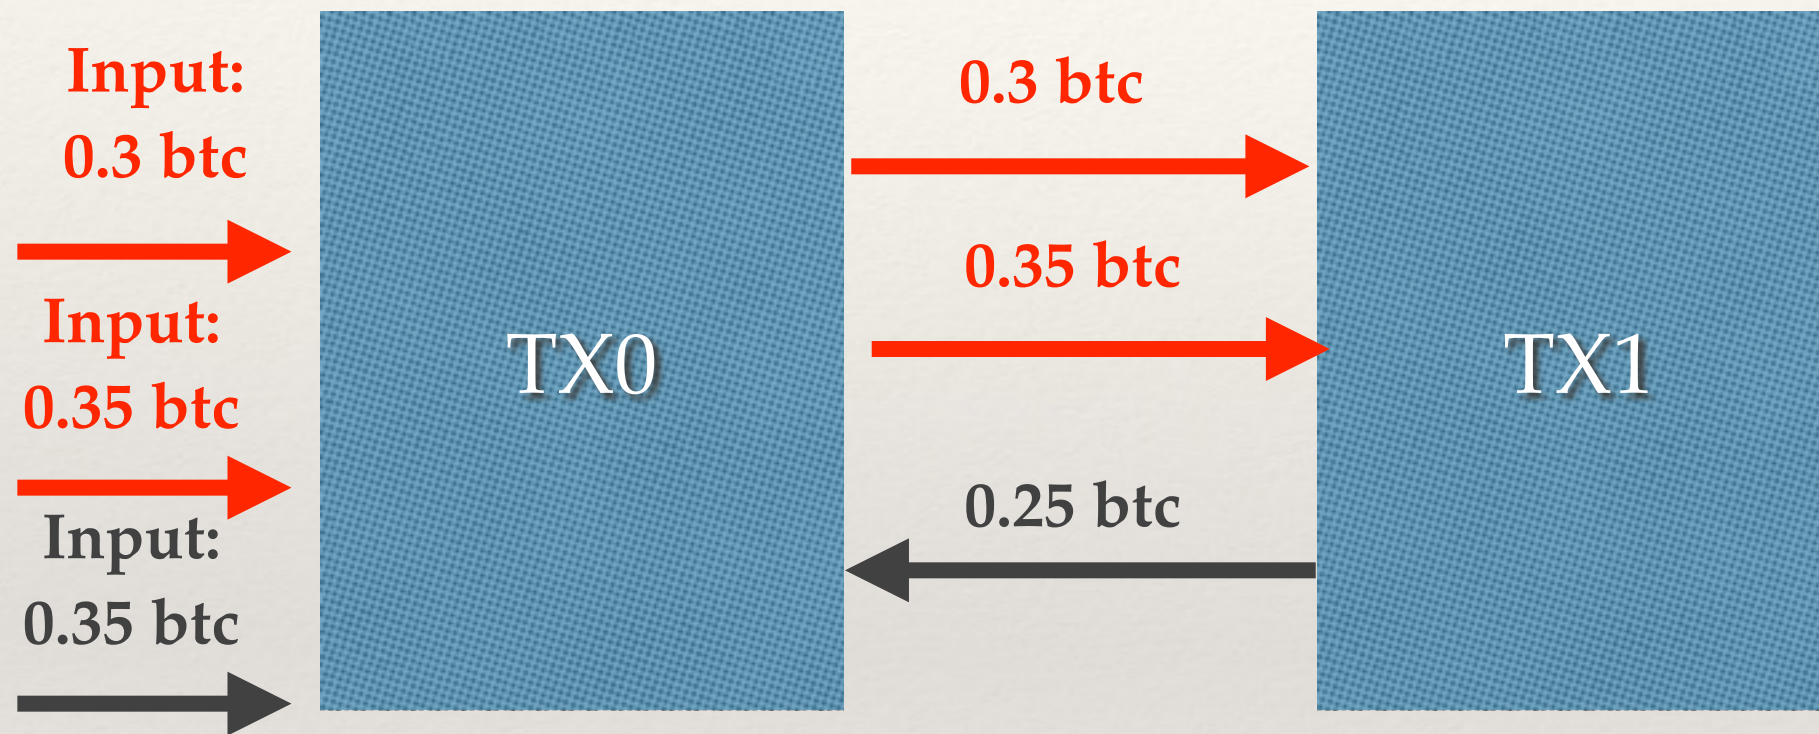
Multiple TX Input/Output

Sending 0.4 btc



Spent vs Unspent Tx's

Sending 0.4 btc



Generator of TX0 now has two UTXOs: 0.35 btc and 0.25 btc

Most bitcoin wallets will just show you $0.35 + 0.25 = 0.6$ btc


$$1 \text{ btc} - 0.4 \text{ btc} = 0.6 \text{ btc}$$

Transaction-Based Ledger

BLOCK	INPUTS/OUTPUTS	SIGNATORY
1	<i>INPUTS:</i> \emptyset <i>OUTPUTS:</i> 25 -> ALICE	N/A
2	<i>INPUTS:</i> 1[0] <i>OUTPUTS:</i> 10 -> BOB; 15 -> ALICE	ALICE
3	<i>INPUTS:</i> 2[1] <i>OUTPUTS:</i> 15 -> CAROL; 10 -> ALICE	ALICE
4	<i>INPUTS:</i> 2[0] <i>OUTPUTS:</i> 10 -> DAN; 8 -> BOB	BOB

Verification Only Back Until Coin Creation

BLOCK	INPUTS/OUTPUTS	SIGNATORY
1	<i>INPUTS:</i> \emptyset <i>OUTPUTS:</i> 25 -> ALICE	N/A
2	<i>INPUTS:</i> 1[0] <i>OUTPUTS:</i> 10 -> BOB; 15 -> ALICE	ALICE
3	<i>INPUTS:</i> 2[1] <i>OUTPUTS:</i> 5 -> CAROL; 10 -> ALICE	ALICE
4	<i>INPUTS:</i> 2[0] <i>OUTPUTS:</i> 10 -> DAN; 8 -> BOB	BOB



Outputs Must Be Entirely Consumed

BLOCK	INPUTS/OUTPUTS	SIGNATORY
1	<i>INPUTS:</i> \emptyset <i>OUTPUTS:</i> 25 -> ALICE	N/A
2	<i>INPUTS:</i> 1[0] <i>OUTPUTS:</i> 10 -> BOB; 15 -> ALICE	ALICE
3	<i>INPUTS:</i> 2[1] <i>OUTPUTS:</i> 15 -> CAROL; 10 -> ALICE	ALICE
4	<i>INPUTS:</i> 2[0] <i>OUTPUTS:</i> 10 -> DAN; 8 -> BOB	BOB

The diagram illustrates the flow of outputs from block 1 to subsequent blocks. A vertical red line originates from the output '25 -> ALICE' of block 1. This line branches into three horizontal red arrows pointing to the left. The top arrow points to the input '1[0]' of block 2. The middle arrow points to the input '2[1]' of block 3. The bottom arrow points to the input '2[0]' of block 4. This visualizes that the entire output of block 1 is consumed by the inputs of blocks 2, 3, and 4.

Joint Payment - 10 Bitcoin To Carol From Bob & Alice

BLOCK	INPUTS/OUTPUTS	SIGNATORY
1	<i>INPUTS:</i> \emptyset <i>OUTPUTS:</i> 25 -> ALICE	N/A
2	<i>INPUTS:</i> 1[0] <i>OUTPUTS:</i> 10 -> BOB; 15 -> ALICE	ALICE
3	<i>INPUTS:</i> 2[0], 2[1] <i>OUTPUTS:</i> 10 -> CAROL; 10 -> ALICE; 5 -> BOB	ALICE, BOB

A Look at a Raw Block

- ❖ That was conceptual - let's look at an actual block (in JSON - key: value format)
- ❖ Blocks generally indicated by hash or height
- ❖ Latest Block (as of Sunday, 3:38 PM EST)
000000000000000000000000f462731f00e4e600469893990c896dba86fa5fb0c1990
- ❖ <https://blockchain.info/rawblock/000000000000000000000000f462731f00e4e600469893990c896dba86fa5fb0c1990>

Anatomy of a Block

- ❖ NOTE: The textbook's description of blocks is slightly outdated, in at least two ways:
 - ❖ Block metadata is slightly off due to SegWit
 - ❖ Does not take into account BIP-34 changes
- ❖ Information in these slides is correct as of the slides' creation date! Bitcoin block structure is liable to change even more in the future.

Anatomy of a Block

- ❖ **Metadata**

- ❖ *Information about the block (size, block number, Merkle root, nonce, etc.)*

- ❖ **Transactions**

- ❖ *Array of all transactions in this block*
 - ❖ *Number of transaction must be ≥ 1 ! (coinbase transaction always present)*

Block Metadata

```
"hash": "000000000000000000000000f462731f00e4e600469893990c896dba86fa5fb0c1990",  
"ver": 536870912,  
"prev_block": "000000000000000000000005f3c0b802ac8b85488d4f88d7a60a20ee4ea3984c1394",  
"mrkl_root": "a65164b813a723570c7d65fad61987036a566d341a0efce4872836a72cb0b934",  
"time": 1537728446,  
"bits": 388454943,  
"fee": 5310126,  
"nonce": 2278125992,  
"n_tx": 1236,  
"size": 498546,  
"block_index": 1724700,  
"main_chain": true,  
"height": 542728,  
"received_time": 1537728446,  
"relayed_by": "0.0.0.0",
```

Compare to Genesis Block

Missing received_time and relayed_by attributes , added later...

```
"hash": "000000000019d6689c085ae165831e934ff763ae46a2a6c172b3f1b60a8ce26f",  
"ver": 1,  
"prev_block": "0000000000000000000000000000000000000000000000000000000000000000",  
"mrkl_root": "4a5e1e4baab89f3a32518a88c31bc87f618f76673e2cc77ab2127b7afdeda33b",  
"time": 1231006505,  
"bits": 486604799,  
"fee": 0,  
"nonce": 2083236893,  
"n_tx": 1,  
"size": 285,  
"block_index": 14849,  
"main_chain": true,  
"height": 0,
```

<https://blockchain.info/rawblock/000000000019d6689c085ae165831e934ff763ae46a2a6c172b3f1b60a8ce26f>

Transactions

- ❖ **Metadata**

- ❖ *Information about the transaction (hash, size, number of inputs, number of outputs, lock time)*

- ❖ **Inputs**

- ❖ *Previous output (all inputs are previous outputs, except coinbase) to be consumed, signature*

- ❖ **Outputs**

- ❖ *Value of output, a Script script (note: no recipient address directly specified here!)*

Transaction Metadata

```
"weight": 900,  
"time": 1537728260,  
"tx_index": 376130660,  
"vin_sz": 1,  
"hash": "1887e38e63ecfcb1490b4adb21b6410b4  
e60454fe09502fb2aa2503b69f3e3c0",  
"vout_sz": 2,  
"relayed_by": "0.0.0.0",  
"lock_time": 542727,  
"ver": 2  
"size": 225,
```

Transaction Inputs

Where is the signature?

```
"inputs": [{
  "sequence": 4294967294,
  "witness": "",
  "prev_out": {
    "spent": true,
    "tx_index": 376128954,
    "type": 0,
    "addr": "17pGR12CGLzhiMzUr18RZzVoEoaUbT8LTa",
    "value": 3108223,
    "n": 1,
    "script": "76a9144ac12bf46ceb296f3c98d19186088b12cb21162688ac"
  },
  "script": "473044022063bfe6f169acd12f4329e6e63b09c1c79856ee07fb0d1
29f9eeb10eccb1ec350022037c94f02090b27047bc882360b4a832b0eaef21bd4
ce239f5332615c5889fec0012103e78877a7225046ff0f9de99d91d0632f1411f
0ad8e7d4e70622bfcc39b616573"
}],
```

Transaction Outputs

```
"out": [{
  "spent": true,
  "tx_index": 376130660,
  "type": 0,
  "addr": "1CXpe1tQQKpeUc89kUqQC2FuiZc5u5AnTa",
  "value": 1017919,
  "n": 0,
  "script": "76a9147e7da0fb7c6edce44f0e88b8ea6d0d933d4bc75888ac"
}, {
  "spent": true,
  "tx_index": 376130660,
  "type": 0,
  "addr": "15JLH7qkwKNJMJfCYrye96hNKqx8SjPkys",
  "value": 2044878,
  "n": 1,
  "script": "76a9142f27aab77d5e5ebc1853076c29b318c032d9375788ac"
}]
```

Programmable Money

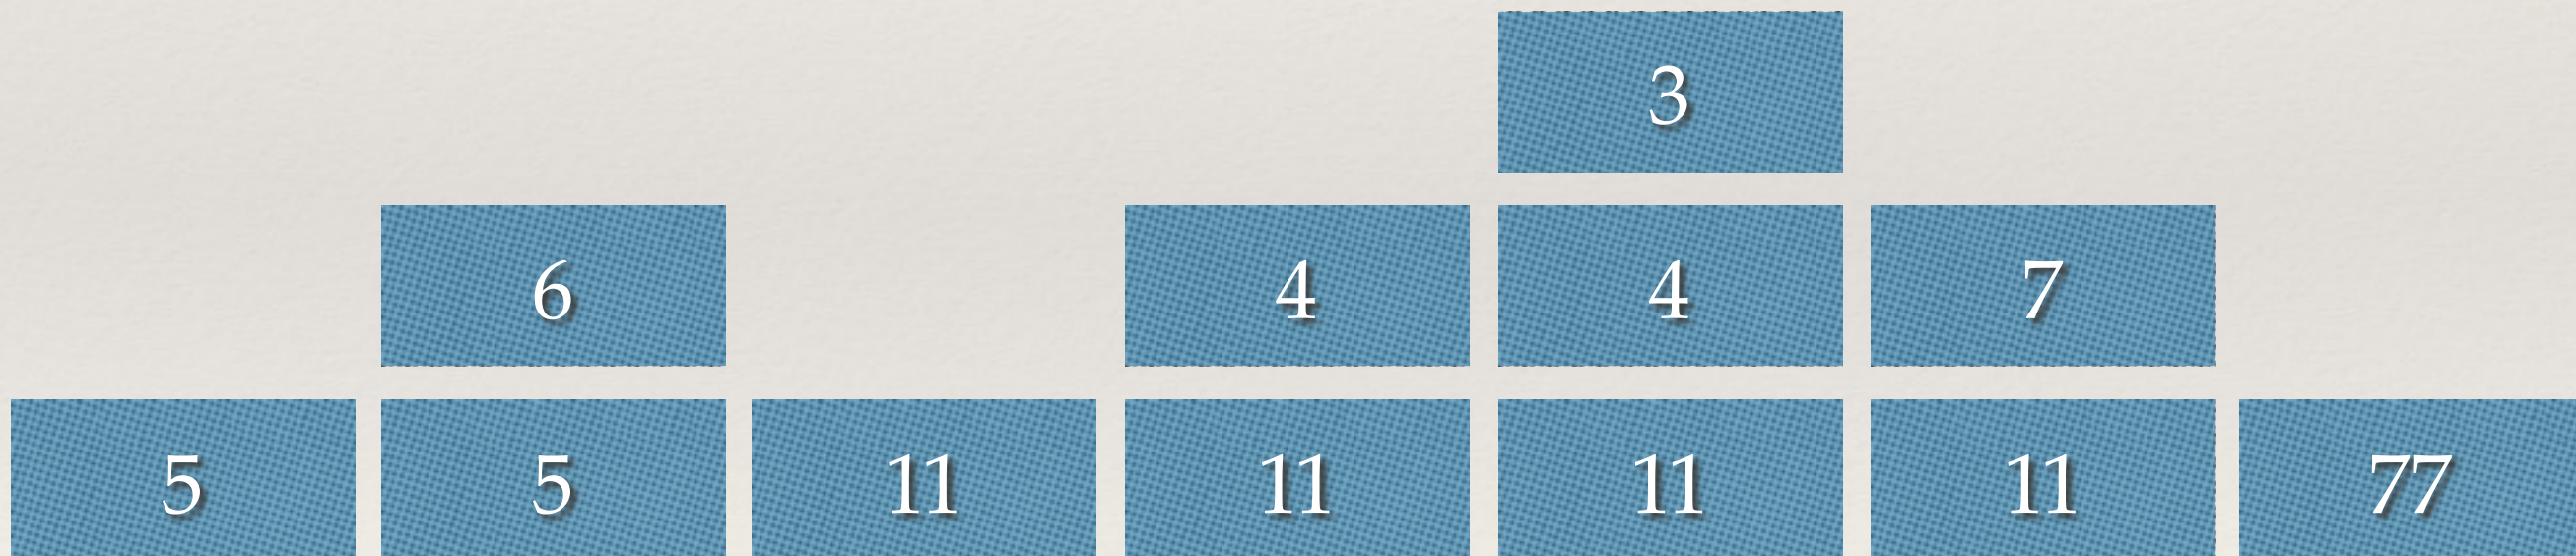
- ❖ Bitcoin is programmable money!
- ❖ A transaction specifies a script to execute
- ❖ The VAST majority of scripts just move previous transaction outputs to a new address, we can (theoretically) do anything that the Script language lets us
- ❖ Although other nodes may have something to say about that...

Script - The Bitcoin Scripting Language

- ❖ Inputs contain scripts (scriptSig attribute) and outputs contain scripts (scriptPubKey) - concatenated and executed
- ❖ Script
 - ❖ Stack-based language
 - ❖ Not Turing-complete (purposely limited)
 - ❖ Native support for complex cryptographic functions
 - ❖ Note: you may see a hex version of this in a blockchain explorer - think of it as the actual bytecode instead of helpful mnemonics

Stack-Based Programming

- ❖ Reverse Polish Notation Calculator
- ❖ $5\ 6\ +\ 4\ 3\ +\ *$



Why Not Turing-Complete?

- ❖ Would never be able to determine if program would end (Halting Problem)
- ❖ Could force arbitrary-length execution on validating nodes (denial-of-service attack)
- ❖ In practice, only a very few kinds of scripts are allowed by most validating nodes (whitelist)
- ❖ See Ethereum (especially concept of *gas*) for one way to have a Turing-complete language running on a blockchain

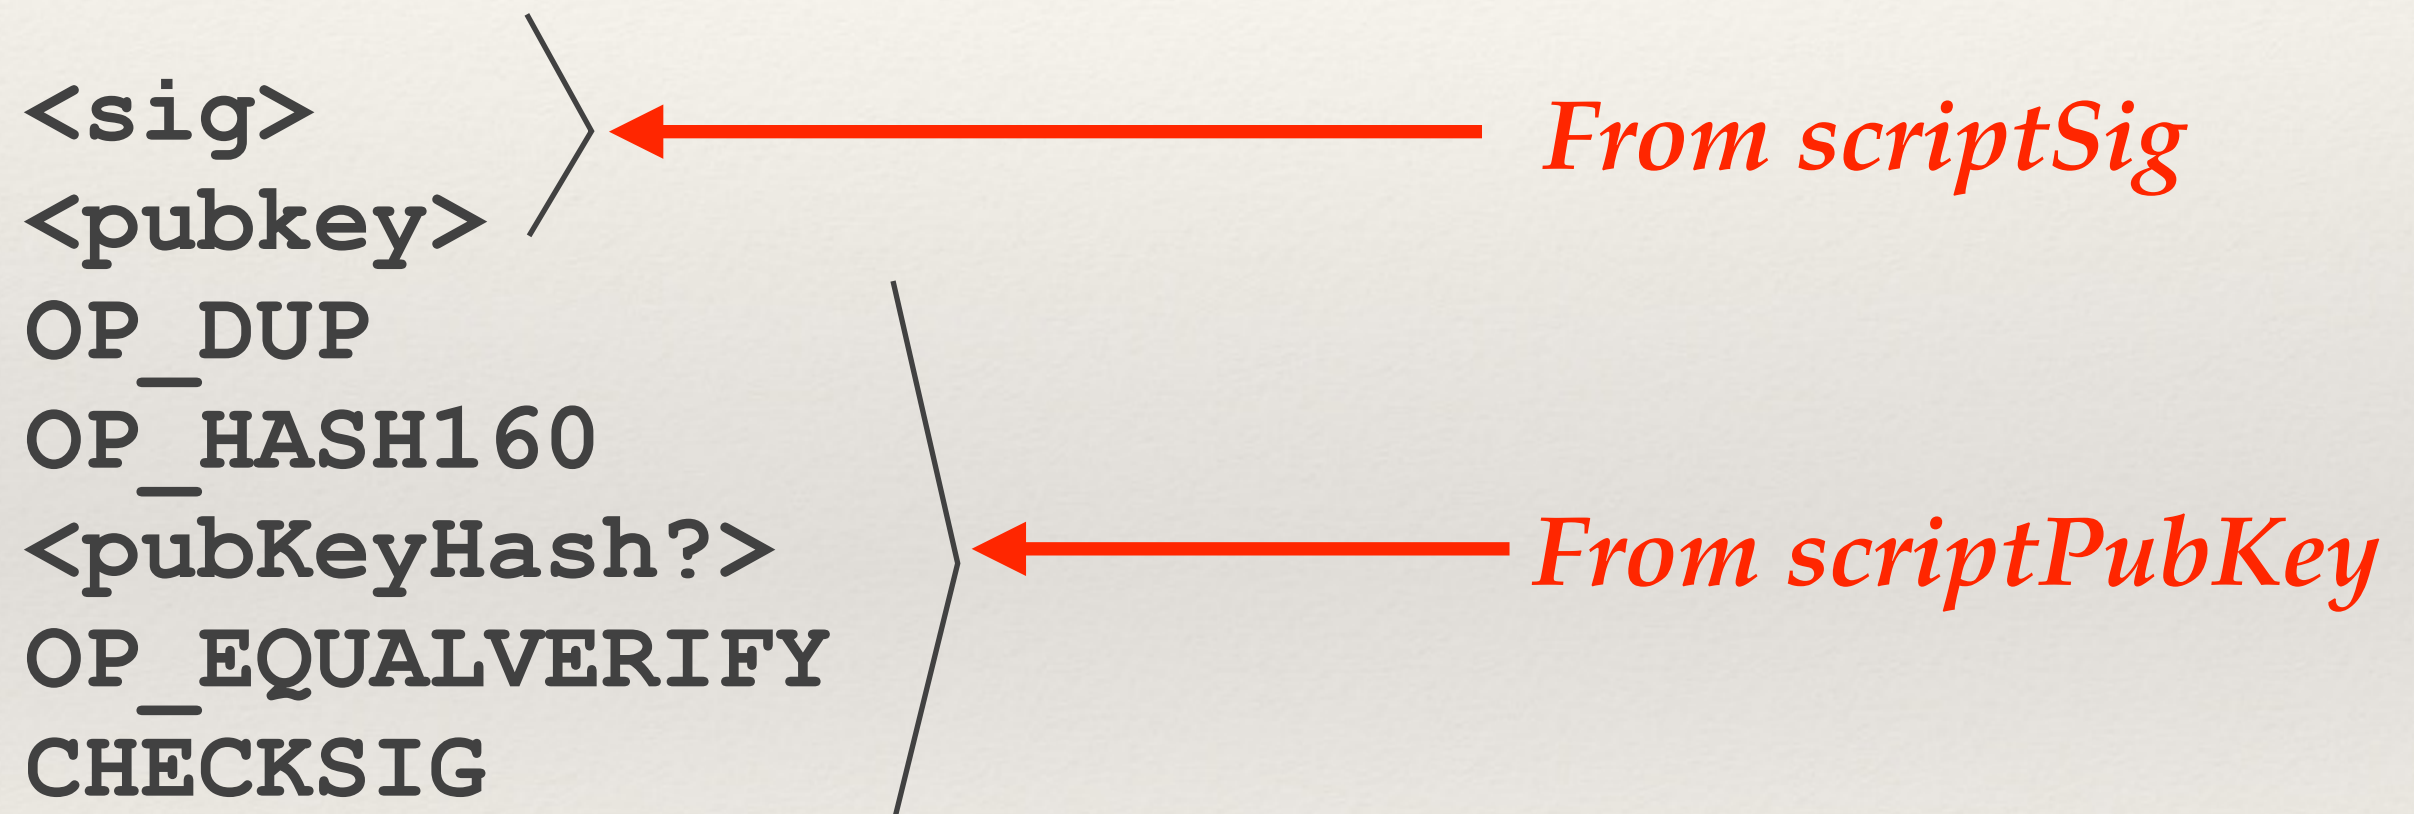
Common Script Commands

- ❖ ***VALUE*** - *Put value on top of stack*
- ❖ **OP_DUP** - *Duplicate top item on stack*
- ❖ **OP_HASH160** - *Hash value on top of stack with SHA-256, then hash again with RIPEMD-160 and put final value back on stack*
- ❖ **OP_EQUALVERIFY** - *Return true if top two values of stack are equal, return false and mark transaction if not*
- ❖ **OP_CHECKSIG** - *Check that the input signature is valid using the input public key for the hash of the current transaction*

Script Execution Walkthrough

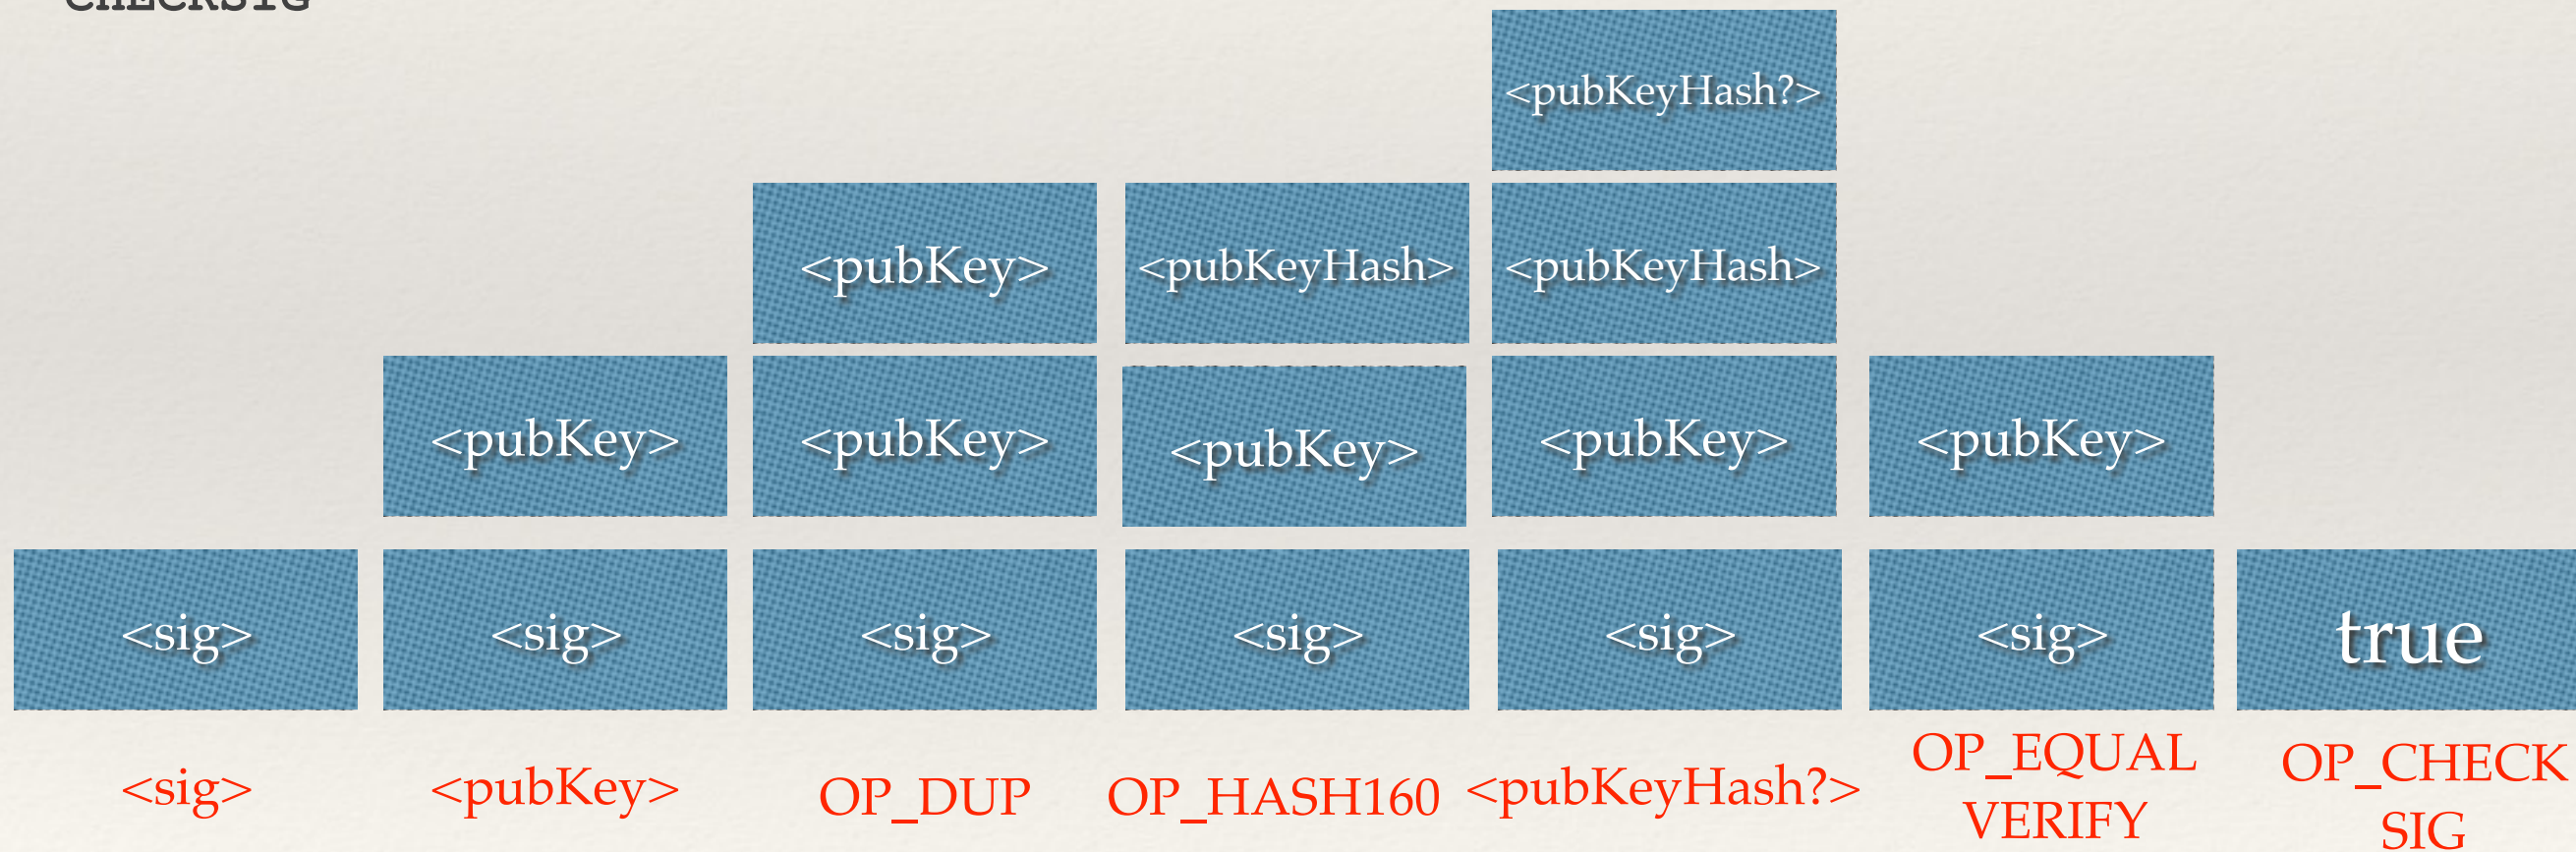
- ❖ Concatenate scriptSig and scriptPubKey
- ❖ Execute script

Script Execution Walkthrough



Basic Bitcoin Script

<sig>
<pubkey>
OP_DUP
OP_HASH160
<pubKeyHash?>
OP_EQUALVERIFY
CHECKSIG



Proof Of Burn

- ❖ Proof of burn establishes that a transaction(s) has been destroyed
- ❖ Why?
 - ❖ Destroy bitcoin to generate some other token
- ❖ How?
 - ❖ Add an OP_RETURN opcode to the scriptPubKey of the transaction - script will always return false

P2PKH vs P2SH

- ❖ P2PKH - *Pay to Public Key Hash* (remember addresses are actually hashes of hashes of a public key, so “pay to address”)
- ❖ P2SH - *Pay to Script Hash* (allow easy multisig or other more complex transactions - not part of original Bitcoin spec, added as BIP-16 in 2012)
- ❖ Additional transaction types:
 - ❖ P2PK - “Pay to Public Key” - generally not used since 0.3, not supported since 0.8
 - ❖ P2WPKH - “Pay To Witness Public Key Hash” - SegWit (Segregated Witness) transaction